

Anchor Properties · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Anchor Properties dbt Wizard demo. Looped unattended on the booth screen.

3-5 days MANUAL MODEL BUILD TODAY	90 sec DBT WIZARD BUILD	\$18M NOI EXPOSURE SIZED	20 hr EARLIER THAN ASSET REVIEW
---	-----------------------------------	------------------------------------	---

ODI architecture · sources, lakehouse, consumers

Fivetran CDC connectors + event streams (Kafka / Kinesis for HL7, telematics, market data)	Iceberg + Glue Open lakehouse on S3, governed catalog, ACID snapshots, time travel	dbt Labs Bronze → Silver → Gold marts, committed as Iceberg snapshots	dbt Wizard Build-time AI · 90s authoring of governed gold marts	Snowflake + 4 Snowflake, Athena, Trino, DuckDB, Spark — same Parquet, no copies
--	--	---	---	---

One set of files on S3. Five engines read the same Parquet through the Glue catalog. Governed end-to-end: SOC 2 · SOX · GAAP · row-level access · column masking.

DATA ENGINEERS ANCHOR HERE Architecture · Wizard · Live · Pipeline Show the loop, the four sub-agents, the materialization, the connector deep-links. This is your DE proof.	BUSINESS USERS ANCHOR HERE Asset Mgrs, Leasing, Property Ops, Investor Relations Show the answers, the KPI tiles, the segments, the executive read. The same governed gold layer powers all of it.
--	--

Frame the demo on business outcomes — never functionality

PAIN Chief Investment Officer can't get answers in time for the Asset Allocation Review	REASON Every new question = 3-5 day DE backlog	IMPACT Missed committee, sized exposure, regulator scrutiny	CAPABILITY Build-time AI on the governed gold layer	PROOF Live model build in 90 seconds — on screen
---	--	---	---	--

The mantra: the win is not that dbt Wizard *can* author a model. The win is that the Chief Investment Officer walks into the next Asset Allocation Review with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties to that one sentence.

The 3-minute script • five beats

00:00 – 00:25 Beat 1 • The stakeholder hook

WHAT YOU SAY (25S)

"It is Friday morning at Anchor Properties. The Chief Investment Officer needs to know: Why is NOI on our office sub-portfolio in three Sunbelt MSAs running 23% below underwriting despite occupancy holding firm at 89%? Asset Allocation Review is in 20 hours. Today that costs three-to-five days of DE work and a missed meeting."

WHAT YOU SHOW

Open the **Scenario page** (`/scenario`). The Chief Investment Officer's question is the first thing in frame. Let the T-minus counter tick.

OUTCOME → Anchors the audience on a real business question, not a feature.

00:25 – 00:50 Beat 2 • The cost of the gap

WHAT YOU SAY (25S)

"There is no gold-layer mart that answers this. Three days from now the team might land it. The Asset Review will have already happened. An \$18 million NOI exposure that has to be sized before the REIT earnings release — and that is just the dollars we can see."

WHAT YOU SHOW

Stay on Scenario. Highlight the four KPI tiles: 3-5 days, 90s, \$18M, T-minus to Asset Allocation Review.

OUTCOME → Quantifies the cost of the gap. Sets up the 90-second proof.

00:50 – 02:20 Beat 3 • The live build (proof moment • DE-anchor)

WHAT YOU SAY (90S)

"We click Run the Build. Four sub-agents — Explorer, Summary, Worker, Verification — open the dbt project, find the right grain, pull a slice of the gold layer, write the new model, write its tests, and materialize it as an Iceberg snapshot. The viewer sees the SQL appear character by character. By the time we are done talking, the model exists, the tests pass, the lineage graph has updated."

Pause once at: "Worker is anchoring the CTE to the grain the Chief Investment Officer asked for — not a notebook-grade approximation."

And again at: "Verification is writing a uniqueness test on the natural key. This is a governed gold asset, not a SQL snippet."

WHAT YOU SHOW

Click into `/wizard-live`. Let the playback run at 1x. The agents type SQL and YAML in real time. The step rail moves Discovery → Schema → Inspection → Model Creation → Tests → Materialization. Terminal-window aesthetic.

OUTCOME → Data engineers see the loop end-to-end. The wow is that the output is a tested, governed asset — not a snippet pasted into Slack.

02:20 – 02:45 Beat 4 · The outcome (P&L tie-back · business-anchor)

WHAT YOU SAY (25S)

"The new mart is materialized. The Asset Allocation Review convenes in 20 hr, not three days. The Chief Investment Officer walks in with the answer: three properties share an elevated TI / leasing-commission run-rate following a recent broker-incentive program; 14 of 19 underperforming buildings share that broker-incentive signature. That is a broker-comp rebalance, not a market repositioning. And the next time anyone asks this — industrial, multifamily, retail — the model is already there."

WHAT YOU SHOW

Cut to /outcome . Materialized model, tests passed, lineage updated, root-cause panel, impact panel. Then quick scroll through the business-user pages (Asset Mgrs, Leasing, Property Ops, Investor Relations) so the audience sees the same gold layer powering the executive read.

OUTCOME → Closes the loop from question to action AND shows the business-user surface running on the same data.

02:45 – 03:00 Beat 5 · The architectural ask

WHAT YOU SAY (15S)

"Fivetran lands every source. Iceberg on S3 holds the bytes; the Glue catalog governs the snapshots. dbt Labs builds the marts. dbt Wizard authors the gaps. Snowflake reads the gold layer — and so do Athena, Trino, DuckDB, and Spark. One set of files, five engines. That loop closes for NOI attribution, lease-rollover risk sizing, ESG disclosure analytics, and SEC + REIT investor reporting — all on the same governed lake. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."

WHAT YOU SHOW

Final card with the architecture loop, QR code to fivetran.com/contact. Tag dbt Labs co-brand.

OUTCOME → Calls out the architectural pattern. Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
Asset Allocation Review gets the answer same-day	3-5 day DE backlog	In time for the meeting, not three days late
\$18M exposure sized in 90 seconds	Open gap window	The cost of the gap closed before lunch
Reusable governed mart for every future ask	One-off notebook scripts	The next person who asks gets it for free
CDO does not add headcount to keep up	Hiring DE backfill	Your team stops being the bottleneck
Auditor sees lineage + tests, not a query history	Manual evidence packs	Examiner-ready by default

Anti-patterns · what NOT to say

Don't say: "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety seconds from question to materialized answer."

Don't say: "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake.

Don't say: "dbt Wizard generates SQL." Say: "dbt Wizard authors a governed asset — schema contract, tests, lineage, ownership."

Don't say: "Look at the typing animation." The animation is the texture, not the value.

Don't open with the architecture diagram. Open with the Chief Investment Officer's question. Architecture is the closer.

Don't skip the business pages. After the wizard finishes, scroll through Asset Mgrs, Leasing, Property Ops, Investor Relations so the audience sees the persona surface that runs on the same gold layer.

Discovery prompts to anchor before the demo

BEFORE YOU PRESS PLAY

- "What questions does your Chief Investment Officer ask that your team cannot answer same-day?"
- "When a stakeholder asks for a new cut of data, what is the average time from ask to merged dbt model?"
- "How many one-off SQL scripts are sitting in someone's Snowflake history that never made it back into the governed layer?"
- "Who is your Economic Buyer for new data investment? Likely audience: CIO, Head of Asset Mgmt, Head of Leasing, Investor Relations."

AFTER THE DEMO LANDS

- "Which of those agent steps — Discovery, Schema, Inspection, Model Creation, Tests, Materialization — is where your team spends the most time today?"
- "If we sized a pilot on three of your highest-value recurring Chief Investment Officer questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Audience: CIO, Head of Asset Mgmt, Head of Leasing, Investor Relations."

Personas at the booth — split the time

DATA ENGINEER · 60% OF BOOTH TIME

Scenario → Live → Outcome → Architecture → Pipeline

Lead with the question, show the four sub-agents, then walk the architecture (Fivetran + Iceberg/Glue + dbt Labs + dbt Wizard + Snowflake/4 engines) and the Pipeline page with real Fivetran connector deep-links.

BUSINESS USER · 40% OF BOOTH TIME

Asset Mgrs, Leasing, Property Ops, Investor Relations

After the wizard finishes, drop into the dashboard pages. Show that the same governed gold layer powers the executive read. The build-time AI is invisible to them — they see the answer.

Required follow-up materials

- One-pager · Anchor Properties · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/real-estate/
- Live demo URL · fivetran-jasonchletsos.github.io/RealEstate-ODI-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the Architecture tab of the demo
- Cross-industry credibility · Pediment Bank

Booth-floor reminder: the demo is built to run unattended on the booth screen. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.

The gold layer gets acted on, not just queried

Activations is the fourth native NewCo stage, right after Transformations. The instant `gold.fct_tenant_credit_risk` flags a tenant with an S&P downgrade or watchlist placement, a lease expiring inside 180 days, and 3%+ of a building's NOI, Activations pushes the renewal-risk segment straight into VTS — no reverse-ETL vendor, no second copy of the data, no analyst in the loop.

TRIGGER 164 days downgrade + expiry <180d + 3%+ NOI	DESTINATION VTS Tenant, Lease, Deal, Task fields	TIME TO OUTREACH 45–60d → same-day manual report cycle, replaced
--	--	--

WHAT YOU SAY	WHAT YOU SHOW
"Here's the part most SEs skip when they explain Activations. We just computed tenant credit risk in the gold layer, <code>fct_tenant_credit_risk</code> — an S&P downgrade, a lease expiring inside 180 days, and a tenant carrying 3% or more of that building's NOI. That's not a dashboard tile, that's a trigger."	Architecture page → Activations panel → trigger condition on <code>fct_tenant_credit_risk</code> .
"Activations reads that segment straight off the governed gold table and pushes it into VTS, the same leasing CRM your reps already live in — no CSV, no ticket to data engineering, no custom script."	/activations-live: Segment Definition → Field Mapping → the five field mappings into VTS .
"Watch the sync run: 342 tenants qualify this morning, and Activations just wrote 12 new high-priority renewal tasks into VTS with the risk score and NOI concentration attached as custom fields on the tenant record."	/activations-live → Sync Preview → API Push → Destination Confirmation table, 342 synced / 12 new tasks .
"That's a 164-day head start on the renewal conversation instead of finding out at the next quarterly portfolio review — protecting the same \$14M in at-risk NOI already flagged on the Investment Committee dashboard."	Click into a VTS task card: "Renewal outreach - high priority - \$2.1M NOI concentration - 164 days to expiry."

Baseline: leasing reps today learn a tenant is a flight risk in the next quarterly manual portfolio review, weeks or months after the S&P rating action posted. Closing that gap is modeled to protect \$14M in at-risk NOI this quarter and cut time-to-outreach from a 45–60 day manual cycle to same-day. Synthetic data — not a real REIT.